

Operating System Structures

Chapter 2

Jobs of the OS

- The OS typically has to provide a few important things
- Support program execution (*processes*)
 - Programs must get a chance to run on a CPU (sometimes)
 - System calls to:
 - Start other programs (e.g., `fork()`, `execl()` on POSIX)
 - Check the status of itself and other programs (e.g., `getpid()`, `getrusage()`, `wait()`)
 - Terminate itself or other programs (e.g., `_exit()`, `kill()`)
- Provide memory for programs
 - Programs need to be able to access memory **without** making a system call every time.
 - System calls to get more memory or give it back (e.g., `sbrk()`, `mmap()`)

Jobs of the OS

- Allocate, use and return files and other resources for programs
 - (e.g., `open()`, `read()`, `close()`)
- Provide protection and security mechanisms
 - (e.g., `chmod()`, `open()`)
- Permit *inter-process communication*
 - If the OS is doing a good job with protection
 - ... then programs are going to need help to work together

Jobs of the OS

- Respond to errors
 - In hardware
 - In user programs
- Perform some accounting
- Provide some kind of user interface
 - Including a collection of *system utilities*
 - Maybe as a *command-line interface*
 - ... or a *graphical interface (or, both)*

Getting Everything Started


- System Boot
- Must get the kernel into memory and start running a few *system utilities / services*.
- At power on, typically:
 - Start running firmware at a known address
 - Load a *boot loader* from secondary storage
 - Bootstrap loader copies kernel into memory and starts executing it
 - Kernel initializes data structures (e.g., interrupt vectors)
 - Kernel starts running at least one system utility (e.g., Unix init or systemd on Linux)

How to Design and Build an OS

- We had to figure out:
 - What an OS should do
 - How it should be designed
- What we'd like
 - Minimize performance overhead
 - Promote portability
 - Promote extensibility
 - Minimize consequences of a failure in the OS
- OS and Hardware have grown up together


Single-Tasking

- Resources scarce and expensive
- Limited hardware support for OS
- MS-DOS Execution Environment
 - One user program at a time
 - Hardware protection not as important



(a)

Waiting for a
Command

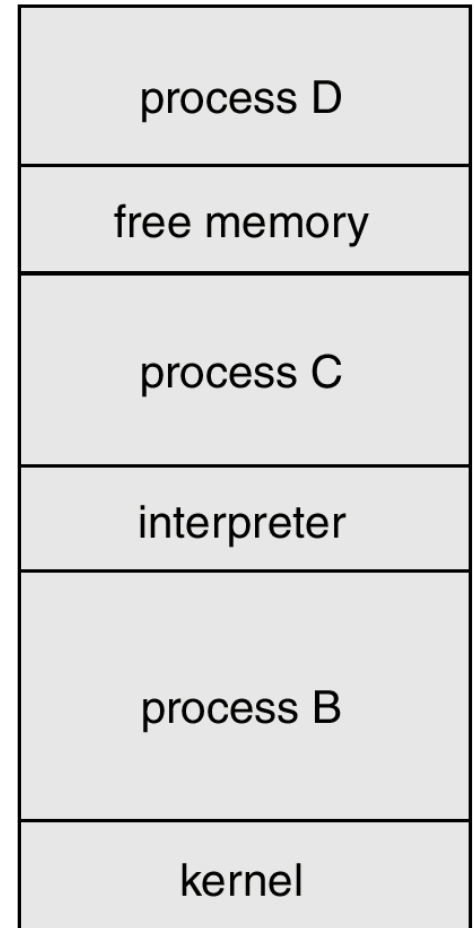


(b)

Running a
Program

Early Multi-Tasking

- With more resources we can:
 - Keep multiple programs (processes) running in memory
 - Let them share resources (including the CPU)
 - That's *multi-tasking*
 - Hardware protection is now important

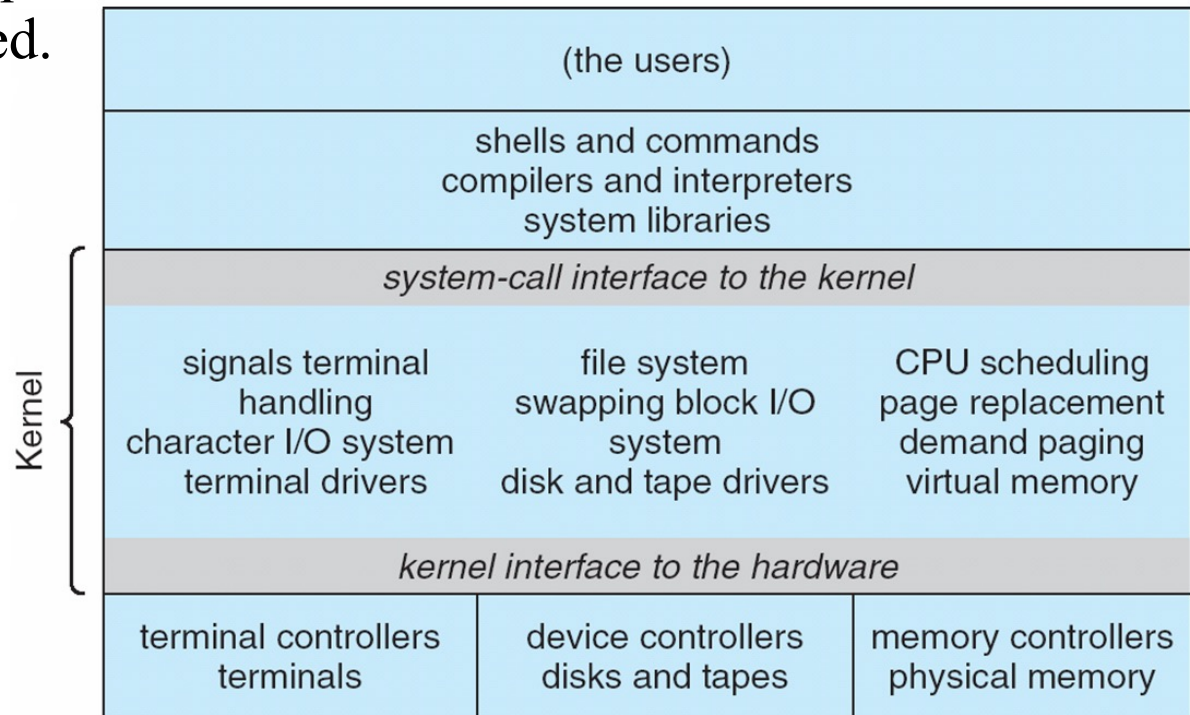


Unix System Structure

- Early OS Organization
 - Limited hardware functionality
 - Little structuring
- Two main components
 - *Kernel*
 - Everything behind the system call interface
 - File System, CPU Scheduling, Memory Management, etc.
 - All running in kernel mode
 - Mostly implemented as a collection of functions
 - *System utilities*, e.g., shell, ls, cat, init, fsck
 - To keep the system running
 - To let users interact with the system.

Unix System Structure

- An example of a *monolithic kernel*
 - All traditional OS services compiled into one big kernel
 - All running in kernel mode.
- This has some disadvantages:
 - Code organization can be unorganized.
 - Harder to modify or extend.
 - Catastrophic if there's a bug in one part of the code.

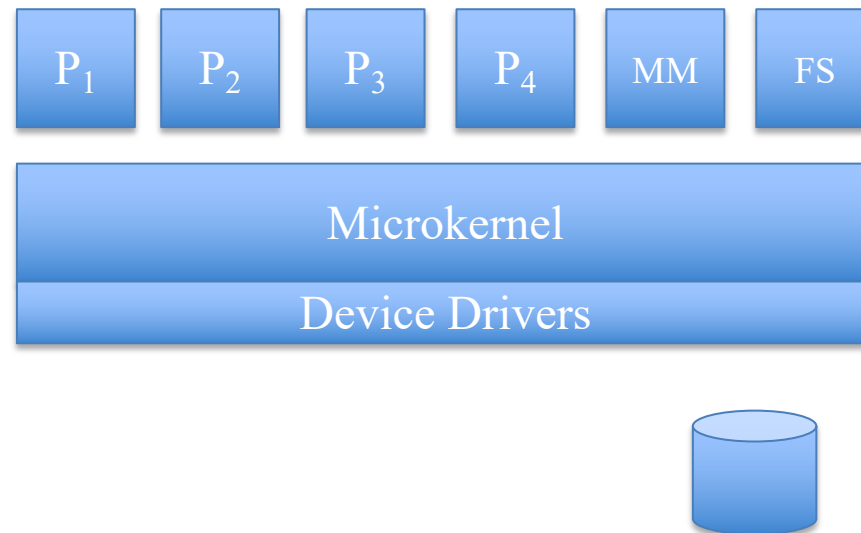


Microkernel Organization


- Move traditional OS components into *user space* (i.e., out of kernel mode)
 - OS functions performed by user-mode processes
 - Do this with as many OS functions as possible
 - e.g., file system
 - So, hopefully, you're left with a very small kernel
- Use of some OS features becomes inter-process communication
- Benefits
 - Easier to extend, add new features in user space
 - Easier to port to a new architecture
 - Less code in kernel mode, so more resistant to failure
 - More secure

Microkernel in Action

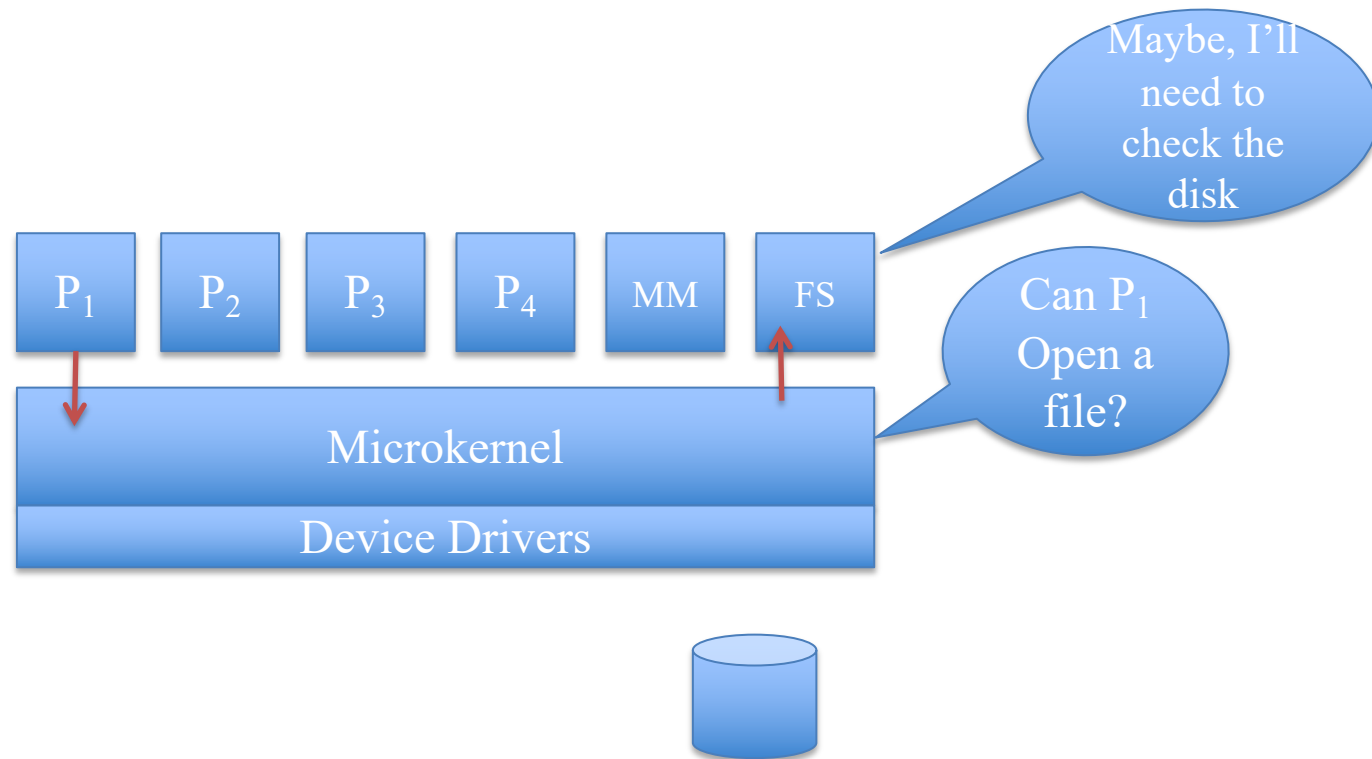
- Pretend user process P_1 needs to open a file



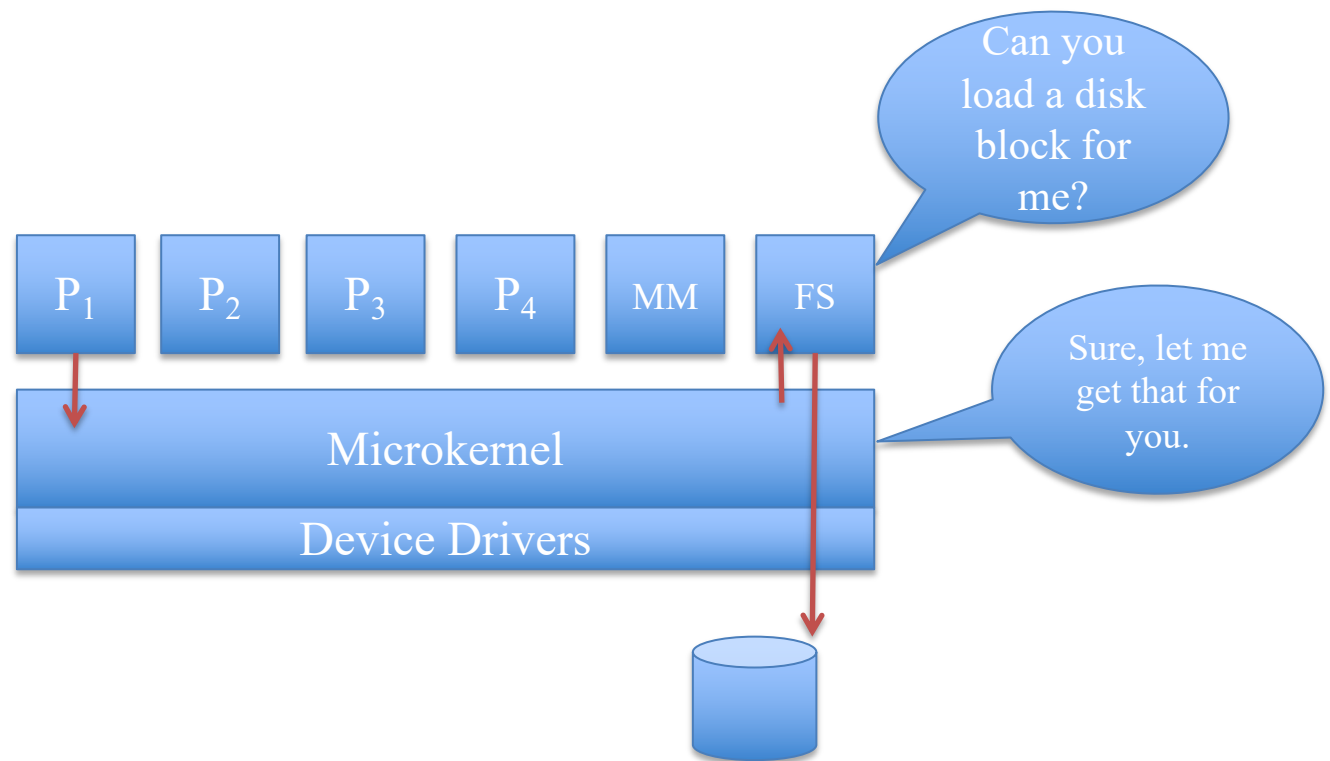
Microkernel in Action




Microkernel in Action



Microkernel in Action



Microkernel in Action



Modular Organization

- Some of the advantages of Microkernel
 - Define categories OS services (e.g., a file system)
 - Each with a well-defined interface
 - e.g., device driver, file system, scheduler
- Kernel *module*
 - Object-oriented approach
 - Each type of module is like an abstract base class
 - Different subclasses, for different module implementations
 - Each module loadable during kernel execution

Microkernel Organization

- Disadvantages
 - More communication & system calls → more overhead
- Requires
 - Well-defined categories of user-mode OS services
 - Well-defined interfaces for each type of service

OS Structures

- Monolithic vs Microkernel
- Almost all modern operating systems are not one pure model, and they use a hybrid approach to combines multiple approaches to address performance, security, usability needs, including Linux, MacOSX, Windows, and Solaris.