

Synchronization

Chapter 6 and 7

Cooperating Threads

- We like our threads to cooperate
- Shared memory can be a nice way to do this
 - For example, shared data structures
 - For example, pool of threads to get work done
- Concurrent access to the same memory
 - Can leave contents in an unknown state
 - Examples, `race1.c`
 - Can be a problem on multiprocessor **and** uniprocessor systems
- We call these *race conditions*
 - Correctness of the program depends on the timing of execution

Race Conditions

- Consider, a single statement may expand to multiple machine instructions
- Execution of these instructions may be interleaved
- How about just `val++` and `val--`

```
load r1, val
inc r1
store val, r1
```

```
load r1, val
dec r1
store val, r1
```

- Interleaving can leave 1 (or -1) in `val`
- This must be happening all the time in our example program!

What's Going On?



```
load r1, val  
inc r1  
store val, r1
```

r1



val



```
load r1, val  
dec r1  
store val, r1
```

r1




What's Going On?



```
load r1, val  
inc r1  
store val, r1
```



```
load r1, val  
dec r1  
store val, r1
```



What's Going On?



```
load r1, val  
inc r1  
store val, r1
```

r1



val




```
load r1, val  
dec r1  
store val, r1
```

r1



What's Going On?

→
load r1, val
inc r1
store val, r1



→
load r1, val
dec r1
store val, r1

What's Going On?

→
load r1, val
inc r1
store val, r1

r1

1

val

0

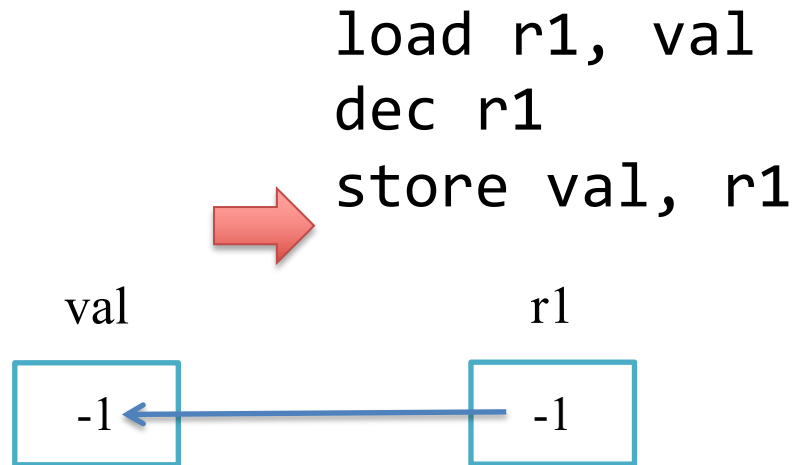

→
load r1, val
dec r1
store val, r1

r1

-1

What's Going On?

→
load r1, val
inc r1
store val, r1




What's Going On?

load r1, val
inc r1
store val, r1



load r1, val
dec r1
store val, r1



Why do we get race conditions?

- OS gets to make CPU scheduling decisions
 - May be difficult to predict what it will do
 - May vary from execution to execution
- Context switch can happen between any two instructions
- Any interleaving of thread execution is possible
- In general, a correct algorithm should:
 - Not depend on choices made by the CPU scheduler
 - Work under concurrent execution
 - i.e., it should be free of race conditions

The Critical Section Problem

- Identify *critical sections* within our code
 - Places where we access shared data
 - Only one thread permitted in any critical section at a time
 - OK, seems like this avoids potential problems
- How about concurrency?
 - Most of the time a thread won't be in a critical section, for typical applications.
 - It will spend most of its time in the *remainder section*

Critical Section Template

```
someThread() {  
    while ( true ) {
```

```
        entry Section;
```

```
        critical Section;
```

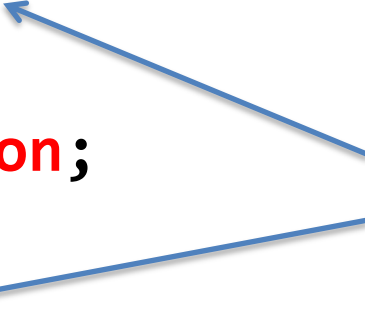
```
        exit Section;
```

```
        remainder Section;
```

```
    }
```

```
}
```

We just need to figure out, once and for all, how to write these two sections so they are guaranteed to work.



Critical Section Problem, let's solve it

```
bool inside[ 2 ] = { false, false };
```

```
someThread0() {  
    while ( true ) {  
  
        while ( inside[ 1 ] ) {  
        }  
        inside[ 0 ] = true;  
  
        critical Section;  
  
        inside[ 0 ] = false;  
  
        remainder Section;  
    }  
}
```

```
someThread1() {  
    while ( true ) {  
  
        while ( inside[ 0 ] ) {  
        }  
        inside[ 1 ] = true;  
  
        critical Section;  
  
        inside[ 1 ] = false;  
  
        remainder Section;  
    }  
}
```


Note, the
body is
empty.




- Does it work? Let's find out

Broken Solution 1

```
bool inside[ 2 ] = { false, false };
```

```
someThread0() {  
    while ( true ) {  
 while ( inside[ 1 ] ) {  
    }  
    inside[ 0 ] = true;  
  
    critical Section;  
  
    inside[ 0 ] = false;  
  
    remainder Section;  
    }  
}
```

```
someThread1() {  
    while ( true ) {  
 while ( inside[ 0 ] ) {  
    }  
    inside[ 1 ] = true;  
  
    critical Section;  
  
    inside[ 1 ] = false;  
  
    remainder Section;  
    }  
}
```

Another Candidate

```
bool wantIn[ 2 ] = { false, false };
```

```
someThread0() {  
    while ( true ) {  
  
        wantIn[ 0 ] = true;  
        while ( wantIn[ 1 ] ) {  
        }
```

critical Section;

```
wantIn[ 0 ] = false;
```

remainder Section;

```
    }  
}
```

```
someThread1() {  
    while ( true ) {  
  
        wantIn[ 1 ] = true;  
        while ( wantIn[ 0 ] ) {  
        }
```

critical Section;

```
wantIn[ 1 ] = false;
```


remainder Section;

```
    }  
}
```

- Does it work?

Another (Broken) Candidate

```
bool wantIn[ 2 ] = { false, false };
```


```
someThread0() {  
    while ( true ) {  
 wantIn[ 0 ] = true;  
    while ( wantIn[ 1 ] ) {  
    }  
}
```

critical Section;

```
wantIn[ 0 ] = false;
```

remainder Section;

```
    }  
}
```

```
someThread1() {  
    while ( true ) {  
 wantIn[ 1 ] = true;  
    while ( wantIn[ 0 ] ) {  
    }  
}
```

critical Section;

```
wantIn[ 1 ] = false;
```

remainder Section;

```
    }  
}
```

Requirements for a Critical-Section Solution

- Mutual Exclusion
 - Only one thread at a time in a critical section
- Progress
 - If two or more threads want to enter the critical section, eventually one will get it.
 - Without needing help from some other thread
 - i.e., threads can't get stuck in the doorway
- Bounded Waiting
 - Can't starve a thread in the entry section
 - A bound on the number of threads that can get ahead of some thread A

Another candidate, take turns

```
int turn = 0;
```

```
someThread0() {  
    while ( true ) {  
  
        while ( turn == 1 ) {  
            }  
  
            critical Section;  
  
            turn = 1;  
  
            remainder Section;  
        }  
    }  
}
```

```
someThread1() {  
    while ( true ) {  
  
        while ( turn == 0 ) {  
            }  
  
            critical Section;  
  
            turn = 0;  
  
            remainder Section;  
        }  
    }  
}
```

- Does it work?

Problems with Taking Turns

```
int turn = 0;
```

```
someThread0() {  
    for ( int i = 0; i < 999;  
          i++ ) {
```

```
    while ( turn == 1 ) {  
    }
```

```
    critical Section;
```

```
    turn = 1;
```

```
    remainder Section;
```

```
    }  
}
```

```
someThread1() {  
    for ( int i = 0; i < 1000;  
          i++ ) {
```

```
    while ( turn == 0 ) {  
    }
```

```
    critical Section;
```

```
    turn = 0;
```

```
    remainder Section;
```

```
    }  
}
```

Peterson's Solution

threads

- Two-~~process~~ solution
- Assume that Load and Store instructions are *atomic*
 - i.e., they can't be interrupted by another instruction
- Use flags, with a turn variable to break ties
 - `bool wantIn[] = { false, false };`
 - `int turn = 0;`
- Set `wantIn[me]` when you want to enter the critical section
- Then, set `turn` to let the other thread go first.
 - Think: “I want in, but you can go first if there's a tie.”

Peterson's Solution

```
bool wantIn[] = { false, false };  
int turn = 0;
```

```
someThread0() {  
    while ( true ) {  
        wantIn[ 0 ] = true;  
        turn = 1; // after you  
        while ( wantIn[ 1 ] &&  
                turn == 1 ) {  
        }  
    }  
}
```

critical Section;

```
wantIn[ 0 ] = false;
```

remainder Section;

```
    }  
}
```

```
someThread1() {  
    while ( true ) {  
        wantIn[ 1 ] = true;  
        turn = 0; // no, after you  
        while ( wantIn[ 0 ] &&  
                turn == 0 ) {  
        }  
    }  
}
```

critical Section;

```
wantIn[ 1 ] = false;
```

remainder Section;

```
    }  
}
```

Beyond Peterson's Solution

- So, we're done, right?
 - Peterson's algorithm solves the critical section problem.
- Not so fast
- What's wrong with just using Peterson's solution as the basis all our synchronization code?

What's Wrong with Peterson's Solution?

- Where should I start?
 - It depends on busy waiting
 - It just works for two threads
 - It's tedious to implement (and we're going to have critical sections all over the place)
 - It's tricky to read and understand (it's not obvious that this is even critical section code)
 - It may not actually work ☹

OS Synchronization Support

- I'd rather have a synchronization mechanism provided by the OS.
 - Then, I wouldn't have to write my own
 - And (more importantly) my process could block when it has to wait.
- We demand synchronization help from the Operating System!

The Semaphore

- Semaphore S, it's like an integer.
 - We'll declare a semaphore like:
`sem s = 1;`
- Two standard operations on a semaphore s.
 - `acquire(s)` and `release(s)`
 - AKA `wait(s)` and `release(s)`
 - AKA `wait(s)` and `post(s)`
 - Originally, `p(s)` and `v(s)`
 - Lots of names for the same two operations . . . not my fault

This is the syntax we'll expect when we write **pseudocode** with semaphores.

Semaphore Operations

- *Acquire* operation : wait until the semaphore has a positive value, then decrement it and proceed
- *Release* operation : increment the semaphore value
- You can think of them this way: (but they aren't really implemented like this)

```
def acquire( sem s ) {  
    while (s <= 0 )  
        ; // no-op  
    s--;  
}  
  
def release( sem s ) {  
    s++;  
}
```

Mutual Exclusion with Semaphores

```
sem s = 1;
```

```
someThread() {  
    while ( true ) {  
  
        acquire( s );  
  
        critical Section;  
  
        release( s );  
  
        remainder Section;  
    }  
}
```

```
someThread() {  
    while ( true ) {  
  
        acquire( s );  
  
        critical Section;  
  
        release( s );  
  
        remainder Section;  
    }  
}
```

Critical Section with Semaphores

- It's a good thing
 - Efficient, we will figure out how to implement `acquire()` using blocking
 - Generalizes to any number of threads/processes
 - Generalizes to any number of critical sections
 - Portable (well, maybe)
 - Ease of use / readability (it looks like synchronization code)
- Plus, semaphores are good at solving lots of synchronization problems
- Let's try some

More than Just Critical Sections

Make sure B() happens after A()

```
someThread() {  
    // Do some stuff here.
```

```
    A();
```

```
    // Do some more stuff.
```

```
}
```

```
otherThread() {  
    // Do some stuff.
```

```
    // Watch TV.
```

```
    B();
```

```
    // Clean room.
```

```
}
```

More than Just Critical Sections

Make sure B() happens after A()

```
sem x = 0;
```

```
someThread() {  
    // Do some stuff here.
```

```
    A();  
    release( x );
```

```
    // Do some more stuff.
```

```
}
```

```
otherThread() {  
    // Do some stuff.
```

```
    // Watch TV.
```

```
    acquire( x );  
    B();
```

```
    // Clean room.
```

```
}
```

A Broken Solution

Make sure B() happens after A()

```
sem x = 1;
```

```
someThread() {  
    // Do some stuff here.
```

```
    A();  
    release( x );
```

```
    // Do some more stuff.
```

```
}
```

```
otherThread() {  
    // Do some stuff.
```

```
    // Watch TV.
```

```
    acquire( x );  
    B();
```

```
    // Clean room.
```

```
}
```


More than Just Critical Sections

Make sure B() happens after A(), and C() happens after B()

someThread() {	otherThread() {	otherOtherThread() {
A();	B();	C();
}	}	}

More than Just Critical Sections

Make sure B() happens after A(), and C() happens after B()

```
sem x = 0;
```

```
sem y = 0;
```

```
someThread() {
```

```
    A();
```

```
    release(x);
```

```
}
```

```
otherThread() {
```

```
    acquire(x);
```

```
    B();
```

```
    release(y);
```

```
}
```

```
otherOtherThread() {
```

```
    acquire(y);
```

```
    C();
```

```
}
```

A Broken Solution

Make sure B() happens after A(), and C() happens after B()

```
sem x = 0;
```

```
someThread() {
```

```
    A();  
    release(x);
```

```
}
```

```
otherThread() {
```

```
    acquire(x);  
    B();  
    release(x);
```

```
}
```

```
otherOtherThread() {
```

```
    acquire(x);  
    C();
```

```
}
```

More than Just Critical Sections

Make sure B() and C() happen after A()

someThread() {	otherThread() {	otherOtherThread() {
A();	B();	C();
}	}	}

Same Solution as Before

Make sure B() and C() happen after A()

```
sem x = 0;
```

```
sem y = 0;
```

```
someThread() {
```

```
    A();
```

```
    release(x);
```

```
}
```

```
otherThread() {
```

```
    acquire(x);
```

```
    B();
```

```
    release(y);
```

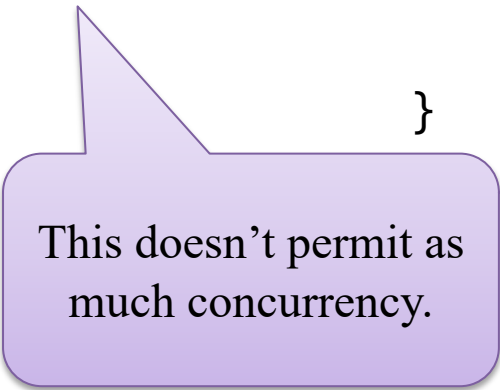
```
}
```

```
otherOtherThread() {
```

```
    acquire(y);
```

```
    C();
```

```
}
```



This doesn't permit as much concurrency.

More than Just Critical Sections

Make sure B() and C() happen after A()

```
sem x = 0;
```

```
sem y = 0;
```

```
someThread() {
```

```
    A();
```

```
    release(x);
```

```
    release(y);
```

```
}
```

```
otherThread() {
```

```
    acquire(x);
```

```
    B();
```

```
}
```

```
otherOtherThread() {
```

```
    acquire(y);
```

```
    C();
```

```
}
```

More than Just Critical Sections

Make sure A() happens after both B() and C()

someThread() {	otherThread() {	otherOtherThread() {
A();	B();	C();
}	}	}

More than Just Critical Sections

Make sure A() happens after both B() and C()

```
sem x = 0;
```

```
sem y = 0;
```

```
someThread() {
```

```
    acquire(x);
```

```
    acquire(y);
```

```
    A();
```

```
}
```

```
otherThread() {
```

```
    B();
```

```
    release(x);
```

```
}
```

```
otherOtherThread() {
```


```
    C();
```

```
    release(y);
```

```
}
```


Bounded-Buffer Problem

- Bounded buffer, a classic synchronization problem
 - Buffer with limited capacity (number of slots)
 - Producer waits if buffer is full
 - Consumer waits if buffer is empty
 - Queue order within the buffer



Implementing Bounded Buffer



```
#define BUFFER_SIZE 10
```

```
// A circular queue of items in the buffer.
```

```
Item buffer[ BUFFER_SIZE ];
```

```
int first = 0;           // Position of first item in the buffer
```

```
int num = 0;             // Number of items in the buffer.
```



A “Solution”

(If you like busy waiting and race conditions)

```
while ( true ) {  
    Item it = makeSomething();  
    while ( num >= BUFFER_SIZE )  
        ; // just wait for a free spot  
    buffer[ ( first + num ) %  
            BUFFER_SIZE ] = it;  
    num++;  
}
```

Producer Thread

```
while ( true ) {  
    while ( num == 0 )  
        ; // Wait for an item  
    Item it = buffer[ first ];  
    first = ( first + 1 ) % BUFFER_SIZE;  
    num--;  
    consumeItem( it );  
}
```

Consumer Thread

Bounded-Buffer with Semaphores

```
sem emptyCount = BUFFER_SIZE;  
sem fullCount = 0;
```

```
while ( true ) {  
    Item it = makeSomething();  
    acquire( emptyCount );  
    buffer[ ( first + num ) %  
            BUFFER_SIZE ] = it;  
    num++;  
    release( fullCount );  
}
```

Producer Thread

```
while ( true ) {  
    acquire( fullCount );  
    Item it = buffer[ first ];  
    first = ( first + 1 ) % BUFFER_SIZE;  
    num--;  
    release( emptyCount );  
    consumeItem( it );  
}
```

Consumer Thread

Fixed?

Bounded-Buffer Solution

```
sem emptyCount = BUFFER_SIZE;  
sem fullCount = 0;  
sem lock = 1;
```

```
while ( true ) {  
    Item it = makeSomething();  
    acquire( emptyCount );  
    acquire( lock );  
    buffer[ ( first + num ) %  
            BUFFER_SIZE ] = it;  
    num++;  
    release( lock );  
    release( fullCount );  
}
```

Producer Thread

```
while ( true ) {  
    acquire( fullCount );  
    acquire( lock );  
    Item it = buffer[ first ];  
    first = ( first + 1 ) % BUFFER_SIZE;  
    num--;  
    release( lock );  
    release( emptyCount );  
    consumeItem( it );  
}
```

Consumer Thread

Fixed!

Bounded-Buffer Solution Recap

- A Semaphore **lock**
 - For preventing concurrent access to the buffer
 - Initialized to 1
 - Just like semaphore-based critical section solution
- A Semaphore **fullCount**
 - Counts the number of filled buffer slots
 - Initialized to 0
- A Semaphore **emptyCount**
 - Counts the number of empty slots
 - Initialized to the buffer capacity
- Note, we can get a lot of value out semaphore counting behavior

Implementing Semaphores

- Inside the OS, we have to be able to:
 - efficiently implement semaphores
 - ... and other synchronization mechanisms
- We'll need to:
 - Solve the critical section problem ourselves
 - Use this to block processes when they have to wait

DIY Synchronization Solutions

- Uniprocessor solutions
 - Could temporarily disable interrupts
 - If we let processes do this, say goodbye to CPU protection
 - Appropriate for *cooperative multitasking*
- Modern CPUs provide help
 - Instructions indented to simplify synchronization
 - Atomic instructions to
 - Test a memory word and set its value
or
 - Compare and swap the contents of two memory words

Test-and-Set-Based Solution

- **Test-And-Set Instruction**

set condition codes based on memory location (is it zero)
set memory location to non-zero

Atomic


```
int lock = 0; // shared
while ( true ) {
    wait:  testAndSet lock
          branchIfNonZero wait

    // Critical section

    lock = 0;

    // Remainder
}
```

This is called
a *spinlock*



Test-and-Set-Based Solution

lock

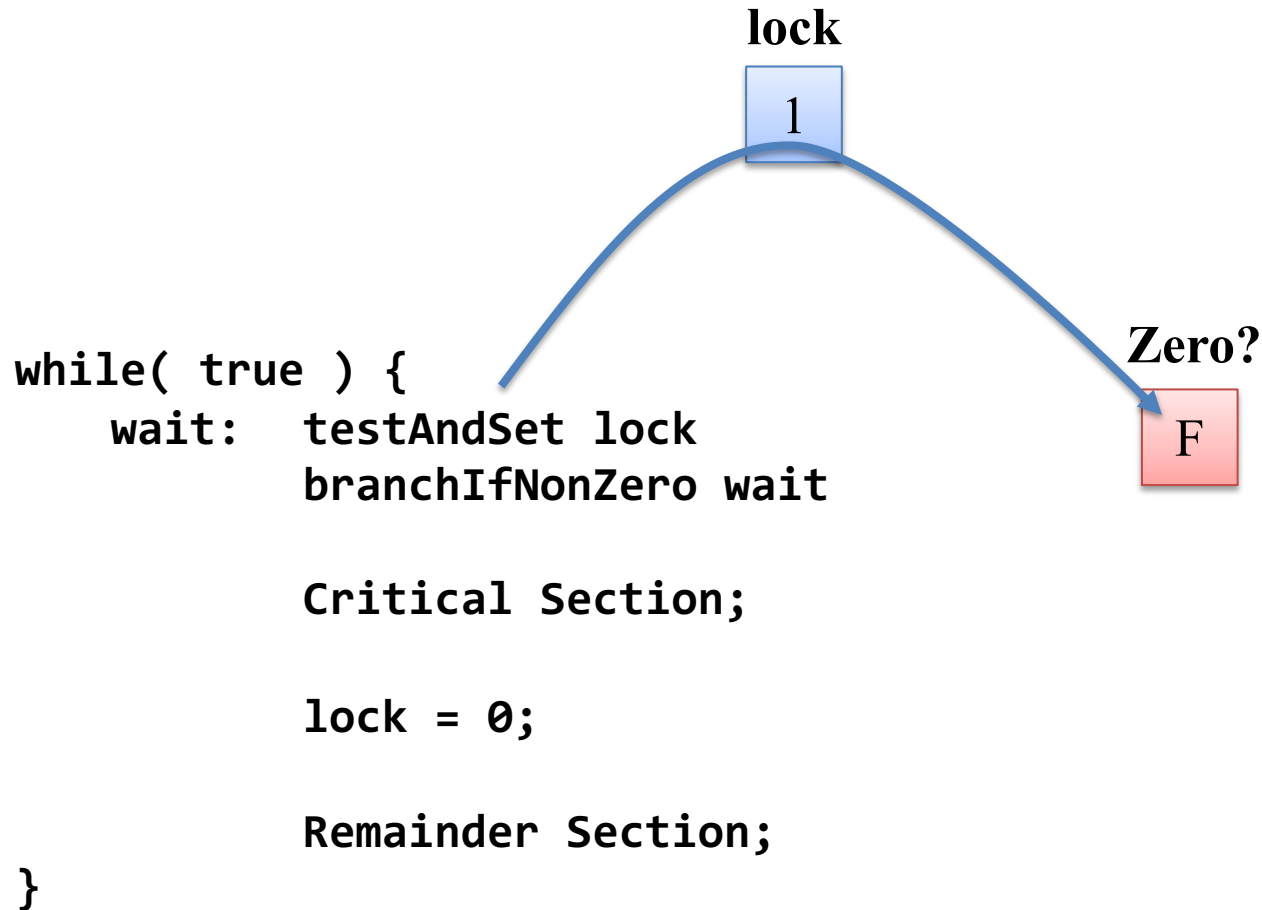


Zero?

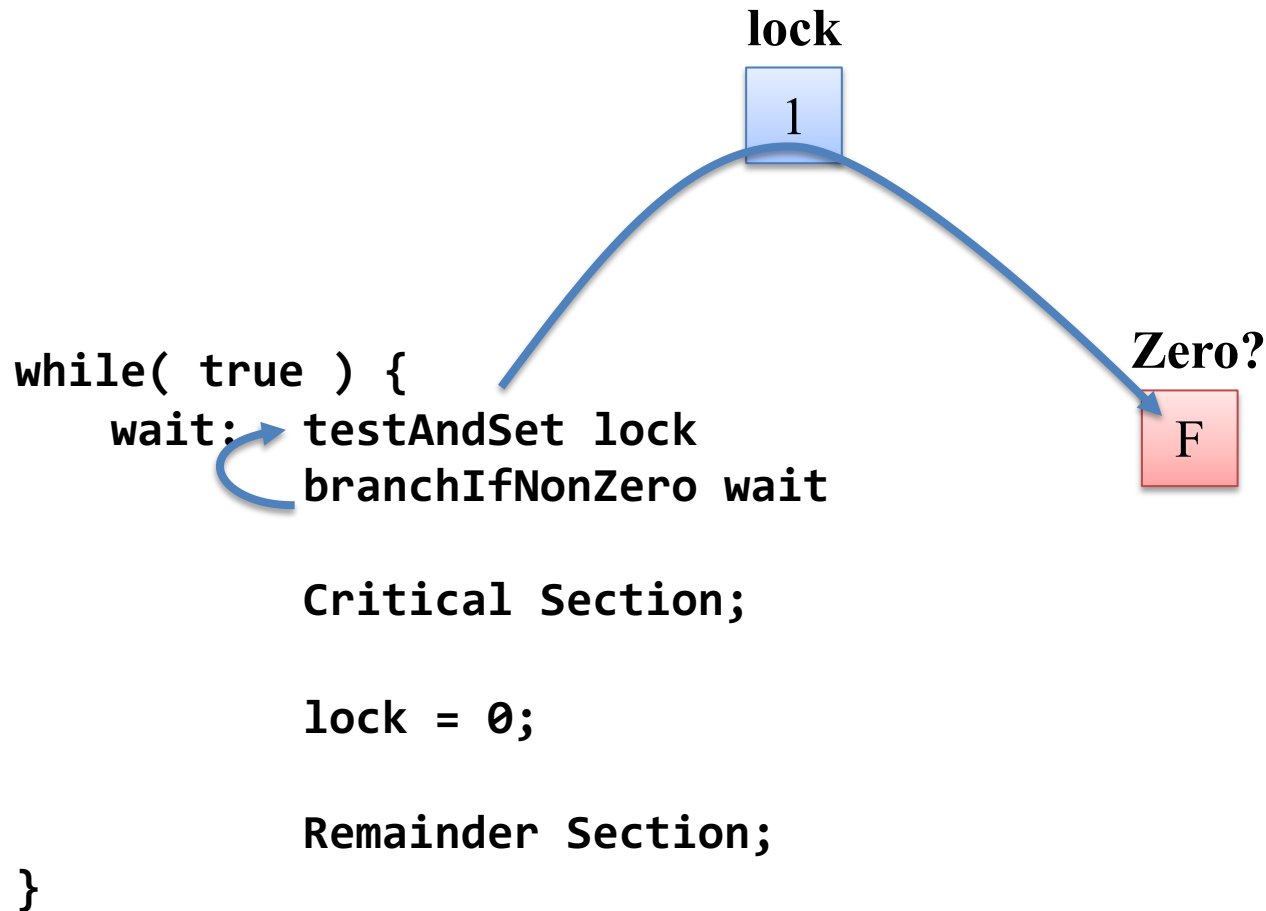


```
while( true ) {  
    wait:  testAndSet lock  
          branchIfNonZero wait  
  
    Critical Section;  
  
    lock = 0;  
  
    Remainder Section;  
}
```

Test-and-Set-Based Solution



Test-and-Set-Based Solution



Test-and-Set-Based Solution

lock

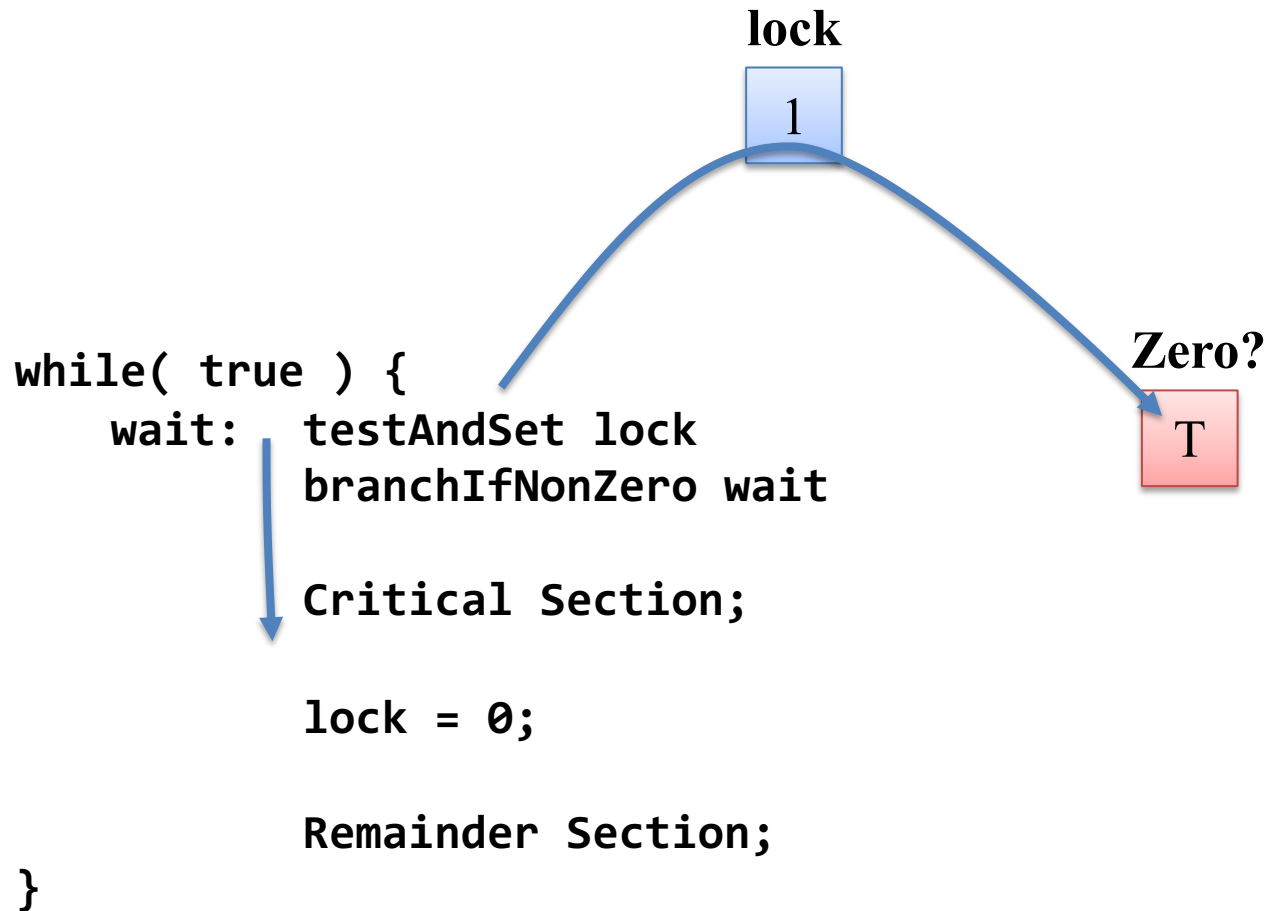


Zero?

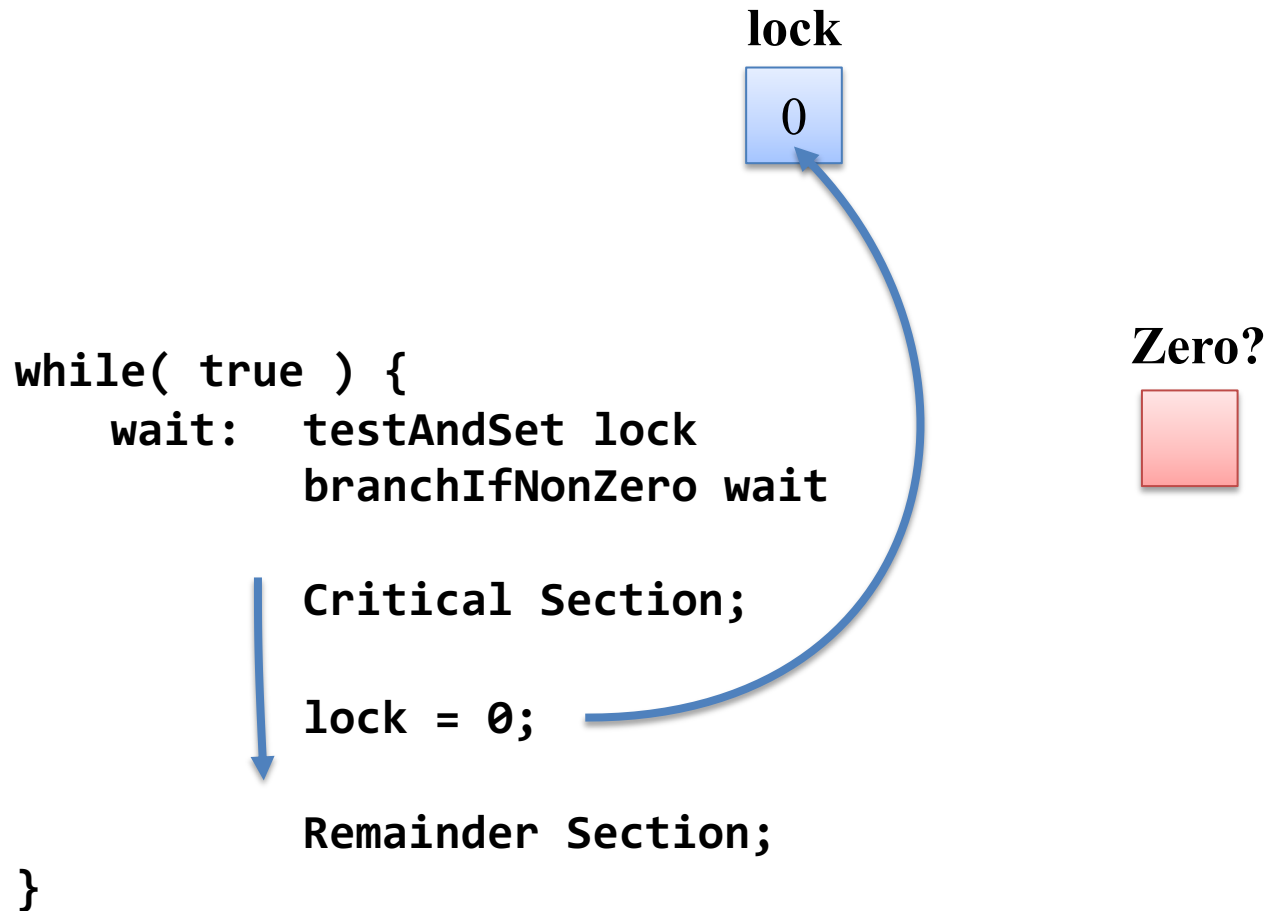


```
while( true ) {  
    wait:  testAndSet lock  
          branchIfNonZero wait  
  
    Critical Section;  
  
    lock = 0;  
  
    Remainder Section;  
}
```

Test-and-Set-Based Solution



Test-and-Set-Based Solution



Compare-and-Swap-Based Solution

- Compare-And-Swap Instruction

```
if ( target == expected ) {  
    set condition codes to true  
    target = newValue;  
} else  
    set condition codes to false
```

Atomic

```
int lock = 0;  // shared  
while ( true ) {  
    wait:  compareAndSwap lock, 0, 1  
           branchIfFalse wait  
  
           // Critical section  
  
           lock = 0;  
  
           // Remainder  
  
}
```


Semaphore Implementation

- So, applications can use semaphores for synchronization
 - If we're the OS, we just need to figure out how to implement `acquire()` and `release()`
 - We could implement them like this, with a busy waiting *spinlock* in each.

```
def acquire( sem s ) {  
    bool proceed = false;  
    while ( !proceed ) {  
        spinlock_entry_code;  
        if ( s > 0 ) {  
            s--;  
            proceed = true;  
        }  
        spinlock_exit_code;  
    }  
}
```

```
def release( sem s ) {  
    spinlock_entry_code;  
    s++;  
    spinlock_exit_code;  
}
```

Semaphore Implementation

- Busy waiting spinlock in `acquire()` and `release()`
 - We know how to do this (and it will work)
 - Don't need to duplicate the spinlock code everywhere
 - All code outside just uses `acquire()` / `release()`
- But, how long might a thread need to wait in `acquire()`?
 - Depends on the application.
- Busy waiting could be a problem

Blocking Semaphore Implementation

- We'd like `acquire()` to just block until the thread can make progress
- We're going to need help from the OS for this
 - The OS can block threads waiting to acquire
 - Keep a queue of PCBs (or threads) for each semaphore, like a device queue
 - So, we have some scheduling queues that aren't associated with any physical device

Blocking Semaphore Implementation

- Pretend the OS has two operations
 - `block()`, mark the current process as blocked
 - `wakeup(P)`, mark process P as ready and place its PCB back on the ready queue

```
acquire(S){  
    spinlock_entry_code;  
    if (s.value <= 0) {  
        enqueue this process on s.list;  
        block();  
        spinlock_exit_code;  
        CPU_Scheduler();  
    } else {  
        s.value--;  
        spinlock_exit_code;  
    }  
}
```

```
release(S){  
    spinlock_entry_code;  
    if (s.value >= 0 && s.list.length > 0) {  
        dequeue a process P from s.list;  
        wakeup(P);  
    } else  
        s.value++;  
    spinlock_exit_code;  
}
```

OS Semaphore Implementation

```
acquire(S){  
    spinlock_entry_code;  
    if (s.value <= 0) {  
        enqueue this process on s.list;  
        block();  
    } else {  
        s.value--;  
    }  
    spinlock_exit_code;  
}
```

```
release(S){  
    spinlock_entry_code;  
    if (s.value >= 0 && s.list.length > 0) {  
        dequeue a process P from s.list;  
        wakeup(P);  
    } else  
        s.value++;  
    spinlock_exit_code;  
}
```

- Busy waiting, but just for a little while
 - Just long enough to inspect the semaphore state ...
 - ... and make a scheduling decision
 - **Doesn't** depend on how long the application needs to wait in acquire().

Flavors of Semaphores

- *Counting Semaphore*
 - Integer value is unrestricted
- *Binary Semaphore*
 - Integer value can only be 0 or 1
 - Good enough, and may be easier to implement
 - Sometimes called a *mutex*
 - Still works fine for critical section solution

```
sem S; // Implicitly set to 1.
```

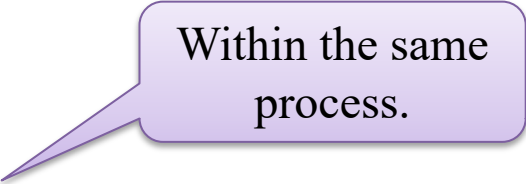
```
acquire( S );  
criticalSection;  
release( S );
```

Do Semaphores Really Work?

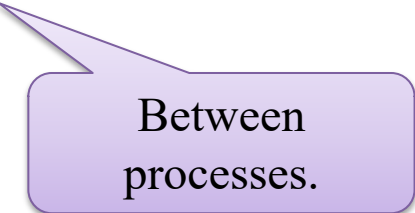
- Yes. They aren't just an idea, they're a readily available synchronization mechanism
- POSIX semaphores
 - Header file: `semaphore.h`
 - Type: `sem_t`
 - `int sem_init(sem_t *sem, int pshared, unsigned int value);`
 - `int sem_destroy(sem_t * sem);`
 - `int sem_wait(sem_t * sem);`
 - `int sem_post(sem_t * sem);`
 - `sem_t * sem_open(char *name, flags, permissions, int value);`
 - `int sem_close(sem_t *sem);`

Do Semaphores Really Work?

- Yes. They aren't just an idea, they're a readily available synchronization mechanism
- POSIX semaphores
 - Header file: semaphore.h
 - Type: sem_t
 - `int sem_init(sem_t *sem, int pshared, unsigned int value);`
 - `int sem_destroy(sem_t * sem);`
 - `int sem_wait(sem_t * sem);`
 - `int sem_post(sem_t * sem);`
 - `sem_t * sem_open(char *name, flags, permissions, int value);`
 - `int sem_close(sem_t *sem);`
- Working examples



Within the same process.




Between processes.

Classic Synchronization Problems

- Some more difficult problems
 - To help us think about synchronization
 - To get some practice using semaphores
 - To see standard types of solutions
 - To improve our computer literacy
- Bounded buffer problem (already seen that)
- Dining-philosophers problem
- Readers-writers problem (later)

Dining-Philosophers Problem

- 5 Philosophers: think or eat
 - Think when you want, eat when you want
 - Each philosopher needs two chopsticks to eat
 - Shared resources: the chopsticks
 - Models cooperating threads that need more than one resource at a time



Semaphore Solution

```
sem chopstick[5] = {1,1,1,1,1}
```

```
philosopher(int i) {  
    while ( true ) {  
        think();
```

repeat


2根

```
        acquire( chopstick[ i ] );  
        acquire( chopstick[(i+1)%5] )  
        eat();
```

grab筷子

```
        release( chopstick[i] );  
        release( chopstick[(i+1)%5] )
```

```
    }  
}
```



Semaphore Solution

- How's my solution
 - Well, it may deadlock
 - How?
- Solutions?
 - Keep an extra chopstick around?
 - Only one at a time can eat? *low efficiency*
 - Only pick up a chopstick if you can get both?
 - Break the symmetry, apply a global ordering to chopstick allocation (this is a reusable trick)

Another POSIX API: Mutex/Lock

- Initialize a mutex lock:
 - `pthread_mutex_t lock;`
 - `pthread_mutex_init(&lock, NULL);`
- Or, we can do it at declaration time:
 - `pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;`
- Acquire the lock:
 - `pthread_mutex_lock(&lock);`
- Release the lock:
 - `pthread_mutex_unlock(&lock);`

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
pthread_mutex_lock(&lock);  
x = x + 1; // or whatever your critical section is  
pthread_mutex_unlock(&lock);
```

Condition Variables (CVs)

- To wait for a condition to become true, a thread can make use of a CV.
- A CV is an explicit queue that threads can put themselves on when some condition is not as desired by **waiting** on the condition.
- Some other thread, when it changes condition, can then wake one (or more) of those waiting threads and thus allow them to continue by **signaling** on the condition.

wait() and signal()

- `wait(cond_t *cv, mutex_t *lock)`
 - assumes the lock is held when `wait()` is called
 - puts caller to sleep + releases the lock (atomically)
 - when awoken, reacquires lock before returning
- `signal(cond_t *cv)`
 - wake a single waiting thread (if ≥ 1 thread is waiting)
 - if there is no waiting thread, just return w/o doing anything

POSIX Condition Variables

- Declare then initialize a condition variable:
 - `pthread_cond_t cond;`
 - `pthread_cond_init(&cond, NULL);`
- Or, do it all at once.
 - `pthread_cond_t cond = PTHREAD_COND_INITIALIZER;`
- Wait/signal on a condition.
 - `pthread_cond_wait(&cond, &lock);`
 - `pthread_cond_signal(&cond);`
- A handy function for waking everyone.
 - `pthread_cond_broadcast(&cond);`

Main thread wants to wait for its spawned thread

```
1 void *child(void *arg) {
2     printf("child\n");
3     // XXX how to indicate we are done?
4     return NULL;
5 }
6
7 int main(int argc, char *argv[]) {
8     printf("parent: begin\n");
9     pthread_t c;
10    pthread_create(&c, NULL, child, NULL);
11    // XXX how to wait for child?
12    printf("parent: end\n");
13    return 0;
14 }
```

Main thread wants to wait for its spawned thread

```
1 volatile int done = 0;
2
3 void *child(void *arg) {
4     printf("child\n");
5     done = 1;
6     return NULL;
7 }
8
9 int main(int argc, char *argv[]) {
10    printf("parent: begin\n");
11    pthread_t c;
12    pthread_create(&c, NULL, child, NULL);
13    while (done == 0)
14        ; // spin
15    printf("parent: end\n");
16    return 0;
17 }
```

Diagram illustrating thread synchronization:

- The child thread sets `done = 1;` (highlighted in a red box) and then calls `thread_exit();` (shown in a red box).
- The main thread enters a `while (done == 0)` loop (highlighted in a red box) to wait for the child thread to finish, and then calls `thread_join();` (shown in a red box).

Implementing Join with CVs (Correct)

```
void thread_exit() {  
    Mutex_lock(&m);  
    done = 1;           // a  
    Cond_signal(&c);     // b  
    Mutex_unlock(&m);  
}
```

```
void thread_join() {  
    Mutex_lock(&m);      // w  
    while (done == 0)    // x  
        Cond_wait(&c, &m); // y  
    Mutex_unlock(&m);    // z  
}
```

Implementing Join with CVs (Wrong 1)

```
void thread_exit() {  
    Mutex_lock(&m);    // a  
    Cond_signal(&c);    // b  
    Mutex_unlock(&m);  // c  
}
```

```
void thread_join() {  
    Mutex_lock(&m);    // x  
    Cond_wait(&c, &m); // y  
    Mutex_unlock(&m);  // z  
}
```

Implementing Join with CVs (Wrong 2)

```
void thread_exit() {  
    done = 1;           // a  
    Cond_signal(&c);     // b  
}
```

```
void thread_join() {  
    if (done == 0)       // x  
        Cond_wait(&c);   // y  
}
```

Implementing Join with CVs (Wrong 3)

```
void thread_exit() {  
    done = 1;           // a  
    Cond_signal(&c);     // b  
}
```

```
void thread_join() {  
    Mutex_lock(&m);      // w  
    if (done == 0)      // x  
        Cond_wait(&c, &m); // y  
    Mutex_unlock(&m);    // z  
}
```

Implementing Join with CVs (Correct?)

```
void thread_exit() {  
    Mutex_lock(&m);  
    done = 1;           // a  
    Cond_signal(&c);     // b  
    Mutex_unlock(&m);  
}
```

```
void thread_join() {  
    Mutex_lock(&m);      // w  
    if (done == 0)      // x  
        Cond_wait(&c, &m); // y  
    Mutex_unlock(&m);    // z  
}
```

Good Rule of Thumb

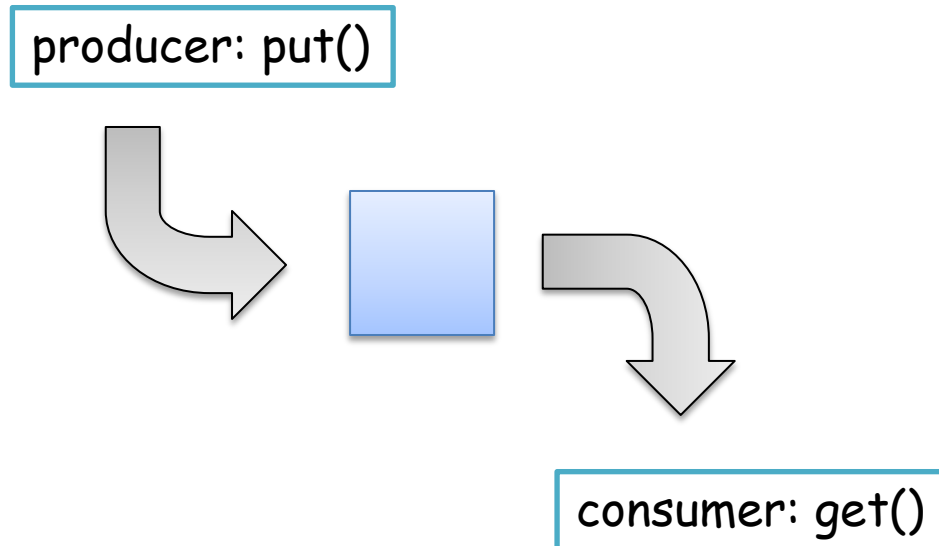
- Keep state in addition to CVs!
 - CVs are used to nudge threads when state changes.
 - If state is already as needed, don't wait for a nudge!
- Always do wait and signal while holding the lock!

Alternative Solution for Producer/Consumer

```
int buffer;  
int count = 0;
```

```
void put(int value) {  
    assert(count == 0);  
    count = 1;  
    buffer = value;  
}
```

```
int get() {  
    assert(count == 1);  
    count = 0;  
    return buffer;  
}
```



Solution V1 (Broken)

```
void *producer(int loops) {  
    for (int i=0; i < loops; i++){  
        Mutex_lock(&m);           //p1  
        if (count == 1)           //p2  
            Cond_wait(&C, &m);    //p3  
        put(i);                   //p4  
        Cond_signal(&C);          //p5  
        Mutex_unlock(&m);         //p6  
    }  
}  
  
void *consumer(int loops) {  
    for (int i=0; i < loops; i++){  
        Mutex_lock(&m);           //c1  
        if (count == 0)           //c2  
            Cond_wait(&C, &m);    //c3  
        int tmp = get();          //c4  
        Cond_signal(&C);          //c5  
        Mutex_unlock(&m);         //c6  
        printf("%d\n", tmp);  
    }  
}
```

Thread Trace: Solution V1 (Broken)

T_{c1}	State	T_{c2}	State	T_p	State	Count	Comment
c1	Running		Ready		Ready	0	
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	Nothing to get
	Sleep		Ready	p1	Running	0	
	Sleep		Ready	p2	Running	0	
	Sleep		Ready	p4	Running	1	Buffer now full
	Ready		Ready	p5	Running	1	T_{c1} awoken
	Ready		Ready	p6	Running	1	
	Ready		Ready	p1	Running	1	
	Ready		Ready	p2	Running	1	
	Ready		Ready	p3	Sleep	1	Buffer full; sleep
	Ready	c1	Running		Sleep	1	T_{c2} sneaks in ...
	Ready	c2	Running		Sleep	1	
	Ready	c4	Running		Sleep	0	... and grabs data
	Ready	c5	Running		Ready	0	T_p awoken
	Ready	c6	Running		Ready	0	
c4	Running		Ready		Ready	0	Oh oh! No data

Solution V2 (Better, But Still Broken)

```
void *producer(int loops) {
    for (int i=0; i < loops; i++){
        Mutex_lock(&m);           //p1
        while (count == 1)       //p2
            Cond_wait(&C, &m);   //p3
        put(i);                  //p4
        Cond_signal(&C);          //p5
        Mutex_unlock(&m);         //p6
    }
}

void *consumer(int loops) {
    for (int i=0; i < loops; i++){
        Mutex_lock(&m);           //c1
        while (count == 0)       //c2
            Cond_wait(&C, &m);   //c3
        int tmp = get();         //c4
        Cond_signal(&C);          //c5
        Mutex_unlock(&m);         //c6
        printf("%d\n", tmp);
    }
}
```

Thread Trace: Solution V2 (Broken)

T_{c1}	State	T_{c2}	State	T_p	State	Count	Comment
c1	Running		Ready		Ready	0	
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	Nothing to get
	Sleep	c1	Running		Ready	0	
	Sleep	c2	Running		Ready	0	
	Sleep	c3	Sleep		Ready	0	Nothing to get
	Sleep		Sleep	p1	Running	0	
	Sleep		Sleep	p2	Running	0	
	Sleep		Sleep	p4	Running	1	Buffer now full
	Ready		Sleep	p5	Running	1	T_{c1} awoken
	Ready		Sleep	p6	Running	1	
	Ready		Sleep	p1	Running	1	
	Ready		Sleep	p2	Running	1	
	Ready		Sleep	p3	Sleep	1	Must sleep (full)
c2	Running		Sleep		Sleep	1	Recheck condition
c4	Running		Sleep		Sleep	0	T_{c1} grabs data
c5	Running		Ready		Sleep	0	Oops! Woke T_{c2}
c6	Running		Ready		Sleep	0	
c1	Running		Ready		Sleep	0	
c2	Running		Ready		Sleep	0	
c3	Sleep		Ready		Sleep	0	Nothing to get
	Sleep	c2	Running		Sleep	0	
	Sleep	c3	Sleep		Sleep	0	Everyone asleep...

Better Solution (usually): use two CVs

Solution V3

```
void *producer(int loops) {
    for (int i=0; i < loops; i++){
        Mutex_lock(&m);           //p1
        while (count == 1)       //p2
            Cond_wait(&E, &m);    //p3
        put(i);                  //p4
        Cond_signal(&F);          //p5
        Mutex_unlock(&m);         //p6
    }
}

void *consumer(int loops) {
    for (int i=0; i < loops; i++){
        Mutex_lock(&m);           //c1
        while (count == 0)       //c2
            Cond_wait(&F, &m);    //c3
        int tmp = get();         //c4
        Cond_signal(&E);          //c5
        Mutex_unlock(&m);         //c6
        printf("%d\n", tmp);
    }
}
```

Summary: rules of thumb


- Keep state in addition to CVs
- Always do wait/signal with lock held
- Whenever you acquire a lock, recheck state

Final Correct Solution for Producer/Consumer

```
int buffer[MAX];  
int fill= 0;  
int use = 0;  
int count = 0
```

```
void put(int value) {  
    buffer[fill] = value;  
    fill = (fill + 1) % max;  
    count ++;  
}
```

```
int get() {  
    int tmp = buffer[use];  
    use = (use + 1) % max;  
    count--;  
    return tmp;  
}
```



Solution V4 (Final)

```
void *producer(void *arg) {
    for (int i=0; i < loops; i++){
        Mutex_lock(&m);           //p1
        while (count == max)     //p2
            Cond_wait(&E, &m);   //p3
        put(i);                  //p4
        Cond_signal(&F);          //p5
        Mutex_unlock(&m);         //p6
    }
}

void *consumer(void *arg) {
    for (int i=0; i < loops; i++){
        Mutex_lock(&m);           //c1
        while (count == 0)       //c2
            Cond_wait(&F, &m);   //c3
        int tmp = get();         //c4
        Cond_signal(&E);          //c5
        Mutex_unlock(&m);         //c6
        printf("%d\n", tmp);
    }
}
```

How to wake the right thread?

- `wait(cond_t *cv, mutex_t *lock)`
 - assumes the lock is held when `wait()` is called
 - puts caller to sleep + releases the lock (atomically)
 - when awoken, reacquires lock before returning
- `signal(cond_t *cv)`
 - wake a single waiting thread (if ≥ 1 thread is waiting)
 - if there is no waiting thread, just return, doing nothing
- `broadcast(cond_t *cv)`
 - wake all waiting threads (if ≥ 1 thread is waiting)
 - if there are no waiting thread, just return, doing nothing

Another Correct Solution V5, But Not As Good As V4

```
void *producer(int loops) {
    for (int i=0; i < loops; i++){
        Mutex_lock(&m);           //p1
        while (count == max)     //p2
            Cond_wait(&C,&m);    //p3
        put(i);                  //p4
        Cond_broadcast(&C);      //p5
        Mutex_unlock(&m);        //p6
    }
}

void *consumer(int loops) {
    for (int i=0; i < loops; i++){
        Mutex_lock(&m);           //c1
        while (count == 0)       //c2
            Cond_wait(&C,&m);    //p3
        int tmp = get();         //c4
        Cond_broadcast(&C);      //c5
        Mutex_unlock(&m);        //c6
        printf("%d\n", tmp);
    }
}
```

Readers-Writers Problem

- Make a distinction for type of concurrent access
 - Two or more threads read concurrently ☺
 - Two or more threads write concurrently ☹
 - A thread writes and some read concurrently ☹
 - Good for controlling access to database fields
 - If we can implement this, we can get more concurrency
- Formally:
 - nw = number of active writers
 - nr = number of active readers
 - Safety condition: $(nr == 0 \text{ and } nw \leq 1) \text{ or } nw == 0$

Reader-Writer Locks

```
acquire_r (rwlock_t *rw) {  
    sem_wait(&rw->lock);  
    rw->readers++;  
    if (rw->readers == 1)  
        sem_wait(&rw->writelock);  
    sem_post(&rw->lock);  
}
```

2

```
typedef struct _rwlock_t {  
    sem_t lock;  
    sem_t writelock;  
    int readers;  
} rwlock_t;  
void rwlock_init(rwlock_t *rw) {  
    rw->readers = 0;  
    sem_init(&rw->lock, 0, 1);  
    sem_init(&rw->writelock, 0, 1);  
}
```

1

```
release_r (rwlock_t *rw) {  
    sem_wait(&rw->lock);  
    rw->readers--;  
    if (rw->readers == 0)  
        sem_post(&rw->writelock);  
    sem_post(&rw->lock);  
}
```

3

```
acquire_w (rwlock_t *rw) {  
    sem_wait (&rw->writelock);  
}  
release_w (rwlock_t *rw) {  
    sem_post (&rw->writelock);  
}
```

4

release lock

What Have We Learned?

- Potential for race conditions
- Critical section problem and DIY solutions
- Synchronization mechanisms
 - Semaphores
 - Mutex + CVs
- Classic problems
 - Bounded-Buffer
 - Dining-Philosophers
 - Readers-Writers